

Algorithms and Data Structures (CSci 115)

California State University Fresno
College of Science and Mathematics
Department of Computer Science
H. Cecotti

Learning outcomes

- More trees
 - B-trees (1971)
 - Definition
 - Search, Insert, Delete an element
- You must know how to implement and use these trees.

B-tree - Definitions

- Generalization of the 2-3 tree

- 2-3 tree: B-tree of order 3.

- Order = Number of children


- B-tree (1971)

- Applications

- Databases
 - Disk-based storage systems

- Features

- 1. keep keys in sorted order for sequential traversing
 - 2. use a hierarchical index to minimize the number of disk reads
 - 3. use partially full blocks to speed insertions and deletions
 - 4. keep the index balanced with a recursive algorithm



B-tree - Definitions

■ Definitions

- All leaves are at same level.
- A B-Tree is defined by the term **minimum degree t**
 - t depends on the “disk block size”.
- **Every node except root** must contain **at least: $t-1$ keys**.
 - Root may contain minimum 1 key.
- All the nodes (with the root) may contain at most: $2t - 1$ keys.
 - A node is **full** if it contains exactly $2t-1$ keys
- Number of children of a node = the number of keys in it **+1**.
- All the keys of a node are sorted in **increasing** order.
 - The child between 2 keys k_1 and k_2 contains all keys in range $]k_1..k_2[$.
- B-Tree grows and shrinks **from root**
 - unlike Binary Search Tree!!
 - BSTs grow **downward** and they shrink from downward.
- Search, insert and delete: $O(\log n)$
 - Like other balanced BSTs

B-tree – 2 Definitions: Degree vs. Order

■ Definitions

➤ Knuth's definition

- A B-tree of **order** m is a tree which satisfies the following properties:
 1. Every node has at most m children.
 2. Every non-leaf node (except the root) has at least $\lceil m/2 \rceil$ child nodes.
 3. The root has at least 2 children if it is not a leaf node.
 4. A non-leaf node with k children contains $k - 1$ keys.
 5. All leaves appear in the same level and carry no information.

➤ → 2-3 Trees: order 3

- “a B-tree of order 3 is a 2-3 tree”

■ Warning

➤ Different definitions and terminologies in the literature!

- Order: 2-3 Trees of order 3
- Degree: 2-3 Trees of degree 2

B-tree - Height

- Height

- The number of disk accesses required for most operations on a B-tree is **proportional** to the **height** of the B-tree
- Theorem

If $n \geq 1$, then for any n -key B-tree T of height h and minimum **degree** $t \geq 2$.


$$h \leq \log_t \frac{n+1}{2}.$$

B-tree – Height

■ Height

➤ Proof

- The root of a B-tree T contains at least 1 key
- All the other nodes contain **at least** $(t-1)$ keys
- $\rightarrow T$ (with height = h) has
 - At least 2 nodes at depth 1
 - At least $2t$ nodes at depth 2
 - At least $2t^2$ nodes at depth 3, ..., until at depth h , it has at least $2t^{h-1}$ nodes.
- \rightarrow number n of **keys** satisfies the inequality:



depth	number of nodes
0	1
1	2
2	$2t$
3	$2t^2$

$$n \geq 1 + (t-1) \sum_{i=1}^h 2t^{i-1} \quad \text{Superior or equal because it is at least } t-1 \text{ keys per node!}$$

$$= 1 + 2(t-1) \left(\frac{t^h - 1}{t - 1} \right)$$

$$= 2t^h - 1. \quad \longrightarrow \quad t^h \leq (n + 1)/2. \quad \longrightarrow \quad h \leq \log_t \frac{n + 1}{2} \quad \square$$

B-tree – Main functions

- Search
- Insert, Delete ...
 - More complicated functions with multiple cases
 - Insert
 - Check a node doesn't get too big
 - Remove
 - Check a node doesn't get too small
- Implementation
 - Remark
 - Previous Trees: simple definition of the node
 - Btree: 2 classes !
 - Btree (main data structure)
 - BTreeNode



B-tree – Main functions

■ Pseudo-code

➤ Search

- With linear search

```
B-TREE-SEARCH( $x, k$ )
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return  $(x, i)$ 
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
```

B-tree – Main functions

- Pseudo-code

- Create an empty B-tree

B-TREE-CREATE(T)

1 $x = \text{ALLOCATE-NODE}()$

2 $x.\text{leaf} = \text{TRUE}$

3 $x.n = 0$

4 $\text{DISK-WRITE}(x)$

5 $T.\text{root} = x$

B-tree – Main functions - Insert

- Insert the new key into an existing leaf node.
- We cannot insert a key into a leaf node that is full !
- → operation **splits** a full node y (having $2t-1$ keys) around its **median key** $y.key_t$ into 2 nodes having only $(t-1)$ keys each.
- The median key moves **up** into y 's parent to identify the dividing point between the 2 new trees.
- But if y 's parent is also full → must split it **before** we can insert the new key
- → we could end up splitting full nodes all the way up the tree !!
- Like a BST, we can insert a key into a B-tree in a single pass **down** the tree from the root to a leaf.
- Not wait to find out whether we will actually need to split a full node in order to do the insertion.
- Instead, travel **down** the tree searching for the position where the new key belongs
- Split each full node we come to along the way (including the leaf itself)
- → whenever we want to split a full node y , we are assured that its parent is not full.



B-tree – Main functions - Insert

- Main pseudo-code

- $O(h)$ disk accesses

B-TREE-INSERT(T, k)

```
1   $r = T.root$ 
2  if  $r.n == 2t - 1$ 
3       $s = \text{ALLOCATE-NODE}()$ 
4       $T.root = s$ 
5       $s.leaf = \text{FALSE}$ 
6       $s.n = 0$ 
7       $s.c_1 = r$ 
8      B-TREE-SPLIT-CHILD( $s, 1$ )
9      B-TREE-INSERT-NONFULL( $s, k$ )
10 else B-TREE-INSERT-NONFULL( $r, k$ )
```

B-tree – Main functions - Insert

■ Insert

1. Initialize x as root.
2. While (x is not leaf) do
 - Find the child of x that is going to be traversed next. Let the child be y.
 - If (y is not full) change x to point to y.
 - If (y is full) split it and change x to point to one of the 2 parts of y.
 - If ($k < \text{mid key}$) in y then set x as 1st part of y
 - Else second part of y.
 - When we **split** y, we move a key from y to its parent x.
3. The loop in step 2 stops when x is leaf.
 - x must have space for 1 extra key! because we have been splitting all nodes in advance.
 - → insert k to x.

B-tree – Main functions - Split

- Input

- a non full internal node x (already allocated in memory)
- an index $i \mid x.c_i$ is a full child of x .

- Procedure

- **splits** this child in 2 and adjusts x so that it has an additional child.
- To split a full root, we will first make the root a child of a new empty root node, so that we can use B-TREE-SPLIT-CHILD.

- → The tree grows in height by one


- **splitting is the only means by which the tree grows.**

B-tree – Main functions - Split

■ Split

B-TREE-SPLIT-CHILD(x, i)

```
1   $z = \text{ALLOCATE-NODE}()$ 
2   $y = x.c_i$ 
3   $z.\text{leaf} = y.\text{leaf}$ 
4   $z.n = t - 1$ 
5  for  $j = 1$  to  $t - 1$ 
6       $z.\text{key}_j = y.\text{key}_{j+t}$ 
7  if not  $y.\text{leaf}$ 
8      for  $j = 1$  to  $t$ 
9           $z.c_j = y.c_{j+t}$ 
10  $y.n = t - 1$ 
11 for  $j = x.n + 1$  downto  $i + 1$ 
12      $x.c_{j+1} = x.c_j$ 
13  $x.c_{i+1} = z$ 
14 for  $j = x.n$  downto  $i$ 
15      $x.\text{key}_{j+1} = x.\text{key}_j$ 
16  $x.\text{key}_i = y.\text{key}_t$ 
17  $x.n = x.n + 1$ 
18 DISK-WRITE( $y$ )
19 DISK-WRITE( $z$ )
20 DISK-WRITE( $x$ )
```




Splitting a node with $t=4$. Node $y=x.c_i$ splits into two nodes, y and z , and the median key S of y moves up into y 's parent

$$2t-1=7$$

B-tree – Main functions - Split

■ Split

➤ Example:



Splitting the root with $t=4$.

- Root node r splits in 2
- A new root node s is created.
 - The new root:
 - a) contains the median key of r
 - b) has the 2 halves of r as children
- The B-tree height ± 1 when the root is split.

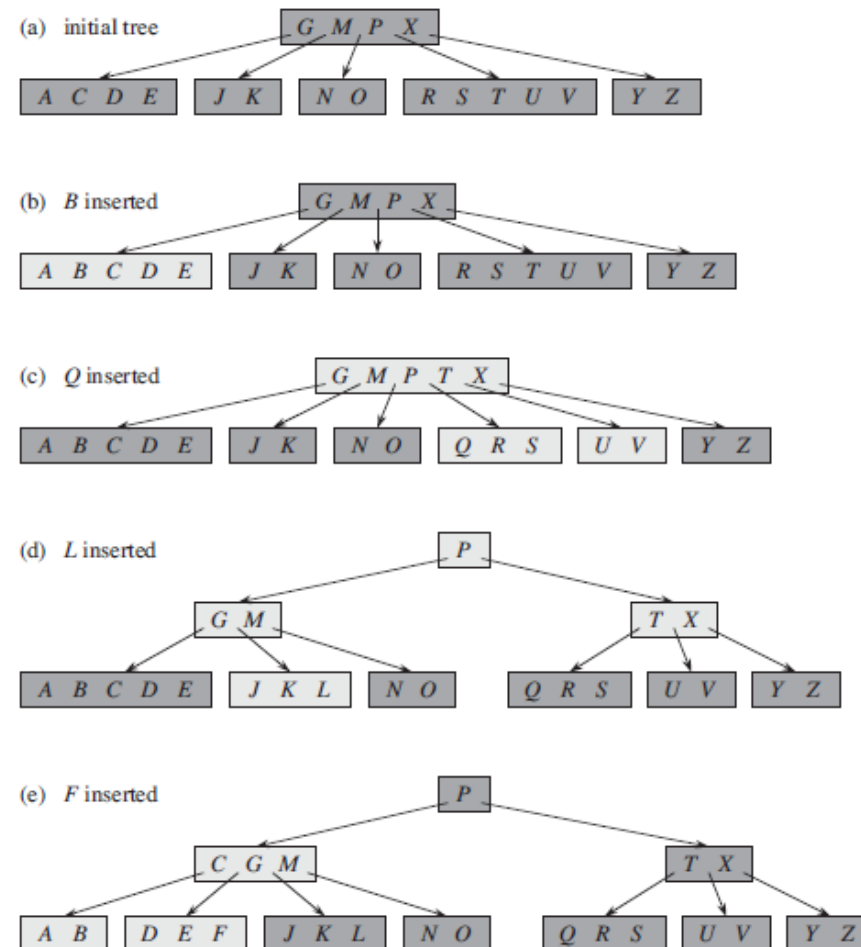
B-tree – Main functions – Insert nonfull

■ Pseudo-code

```
B-TREE-INSERT-NONFULL( $x, k$ )
1   $i = x.n$ 
2  if  $x.leaf$ 
3      while  $i \geq 1$  and  $k < x.key_i$ 
4           $x.key_{i+1} = x.key_i$ 
5           $i = i - 1$ 
6       $x.key_{i+1} = k$ 
7       $x.n = x.n + 1$ 
8      DISK-WRITE( $x$ )
9  else while  $i \geq 1$  and  $k < x.key_i$ 
10      $i = i - 1$ 
11      $i = i + 1$ 
12     DISK-READ( $x.c_i$ )
13     if  $x.c_i.n == 2t - 1$ 
14         B-TREE-SPLIT-CHILD( $x, i$ )
15         if  $k > x.key_i$ 
16              $i = i + 1$ 
17     B-TREE-INSERT-NONFULL( $x.c_i, k$ )
```

B-tree – Main functions - Insert

■ Example:



B-tree – Main functions – Delete

- Delete the key k
 - Constraints
 - Must guard against deletion producing a tree whose structure violates the B-tree properties
 - Check a node doesn't get too small during deletion!

B-tree – Main functions – Delete

- Delete the key k (main function)

1. If k is in node x and x is a leaf \rightarrow delete the key k from x .
2. If it is in node x and x is an internal node, \rightarrow do :
 - a) If the child y that precedes k in node x has at least t keys, then find the predecessor k_0 of k in the sub-tree rooted at y . Recursively delete k_0 , and replace k by k_0 in x .
can find k_0 and delete it in a single downward pass.
 - b) If y has fewer than t keys, then, symmetrically, examine the child z that follows k in node x . If z has at least t keys, then find the successor k_0 of k in the subtree rooted at z . Recursively delete k_0 , and replace k by k_0 in x . (We can find k_0 and delete it in a single downward pass.)
 - c) Else, if both y and z have only $t-1$ keys, merge k and all of z into y , so x loses both k and the pointer to z , y has now contains $2t-1$ keys. Then free z and recursively delete k from y ...
3. If k is not present in internal node x , determine the root $x.c(i)$ of the appropriate subtree that must contain k (if k is in the tree at all)
If $x.c(i)$ has only $(t-1)$ keys, execute step 3a or 3b as necessary to guarantee that we descend to a node containing at least t keys.
 - Then finish by recursing on the appropriate child of x
 - a) If $x.c(i)$ has only $t-1$ keys but has an immediate sibling with at least t keys, give $x.c(i)$ an extra key by moving a key from x down into $x.c(i)$, moving a key from $x.c(i)$'s immediate left or right sibling up into x , and moving the appropriate child pointer from the sibling into $x.c(i)$.
 - b) if $x.c(i)$ and both of $x.c(i)$'s immediate siblings have $t-1$ keys, merge $x.c(i)$ with one sibling, which involves moving a key from x down into the new merged node to become the median key for that node.



B-tree – Main functions – Delete

■ Cases

1. **If** the key k is in node x and x is a leaf **then** delete the key k from x .
2. **If** the key k is in node x and x is an internal node **then**
 - a) **If** the child y that precedes k in node x has at least t keys, then find the predecessor k' of k in the subtree rooted at y .
Recursively delete k' , and replace k by k' in x .
(find k' and delete it in a single downward pass)
 - b) **If** y has fewer than t keys, **then**, symmetrically, examine the child z that follows k in node x .
If z has at least t keys, **then** find the successor k' of k in the subtree rooted at z .
Recursively delete k' , and replace k by k' in x .
(can find k' and delete it in a single downward pass)
 - c) **Else, if** both y and z have only $(t - 1)$ keys, merge k and all of z into y ,
so that x loses both k and the pointer to z , and y contains $2t - 1$ keys.
Then free z , and recursively delete k from y .

B-tree – Main functions – Delete

■ Cases

3. If the key k is not present in internal node x , determine the root $x.c_i$ of the appropriate subtree that must contain k if k is in the tree at all.

If $x.c_i$ has only $(t-1)$ keys, execute step **3a** or **3b** as necessary to guarantee that we descend to a node containing at least t keys.





Then finish by recursing on the appropriate child of x .

- a) If $x.c_i$ has only $(t-1)$ keys but has an immediate sibling with at least t keys
 - give $x.c_i$ an extra key by moving a key from x down into $x.c_i$,
 - moving a key from $x.c_i$'s immediate left or right sibling up into x
 - and moving the appropriate child pointer from the sibling into $x.c_i$.
- b) If $x.c_i$ and both of $x.c_i$'s immediate siblings have $(t-1)$ keys, merge $x.c_i$ with one sibling, which involves moving a key from x down into the new merged node to become the **median** key for that node.

B-tree – Main functions – Delete

■ Example:

- The minimum degree for this B-tree is $t=3$
- → a node (diff than the root) cannot < than 2 keys
 - (modified nodes = lightly shaded)
- Deletion of F:
 - **case 1**: simple deletion from a leaf
- Deletion of M:
 - **case 2a**: the predecessor L of M moves **up** to take M's position.
- Deletion of G:
 - **case 2c**: push G **down** to make node "DEGJK" and then delete G from this leaf.



B-tree – C++ implementation

■ B-tree Class

➤ Only 3 main methods !

```
class BTree {
    BTreeNode *root; // Pointer to root node
    int t;           // Minimum degree
public:
    // Constructor (Initializes tree as empty)
    BTree(int _t) { root = NULL; t = _t; }
    void traverse() { if (root != NULL) root->traverse(); }
    // search a key in this tree
    BTreeNode* search(int k) { return (root == NULL) ? NULL : root->search(k); }
    // insert a new key in the B-Tree
    void insert(int k);
    // remove a key in the B-Tree
    void remove(int k);
};
```


B-tree – C++ implementation

- BTreeNode Class

- Main part of the data structure

```
class BTreeNode {
    int *keys;      // Array of keys
    int t;          // Minimum degree (defines the range for number of keys)
    BTreeNode **C;  // An array of child pointers
    int n;          // Current number of keys
    bool leaf;      // true == node is leaf
public:
    BTreeNode(int _t, bool _leaf); // Constructor
    // traverse all nodes in a subtree rooted with this node
    void traverse();
    // search a key in subtree rooted with this node, returns NULL if k is not present
    BTreeNode *search(int k);
    // returns the index of the first key >= k
    int findKey(int k);
    // insert a new key in the subtree rooted with this node
    // the node must be non-full when insertNonFull is called
    void insertNonFull(int k);
    // split the child y of this node. i= index of y in child array C[]. child y must be full when splitchild is called
    void splitChild(int i, BTreeNode *y);
    // remove the key k in subtree rooted with this node
    void remove(int k);
    // remove the key present in idx-th position in this node which is a leaf
    void removeFromLeaf(int idx);
    // remove the key present in idx-th position in this node which is a non-leaf node
    void removeFromNonLeaf(int idx);
    // get the predecessor of the key present in the idx-th position in the node
    int getPred(int idx);
    // get the successor of the key present in the idx-th position in the node
    int getSucc(int idx);
    // fill up the child node present in the idx-th position in the C[] array if that child has less than t-1 keys
    void fill(int idx);
    // borrow a key from the C[idx-1]-th node and place it in C[idx]th node
    void borrowFromPrev(int idx);
    // borrow a key from the C[idx+1]-th node and place it in C[idx]th node
    void borrowFromNext(int idx);
    // merge idx-th child of the node with idx+1 th child of the node
    void merge(int idx);
    // Set BTree friend of this --> access private members of BTreeNode in BTree methods
    friend class BTree;
};
```



B-tree – C++ implementation

- B tree methods
 - Insert function

```
// The main function that inserts a new key in this B-Tree
void BTree::insert(int k) {
    if (root == NULL) {
        root = new BTreeNode(t, true);
        root->keys[0] = k; // Insert key
        root->n = 1; // Update number of keys in root
    }
    else {
        // If root is full, then tree grows in height
        if (root->n == 2 * t - 1) {
            BTreeNode *s = new BTreeNode(t, false);
            s->C[0] = root; // set old root as child of new root
            // split the old root and move 1 key to the new root
            s->splitChild(0, root);
            // new root has 2 children.
            //Decide which of the 2 children is going to have new key!
            int i = 0;
            if (s->keys[0] < k)
                i++;
            s->C[i]->insertNonFull(k);
            // Change root
            root = s;
        }
        else // root is not full, call insertNonFull for root
            root->insertNonFull(k);
    }
}
```

B-tree – C++ implementation

- B tree methods

- Remove function

```
void BTree::remove(int k) {  
    if (!root) {  
        cout << "The tree is empty.\n";  
        return;  
    }  
    // Call the remove function for root  
    root->remove(k);  
    // if the root node has 0 keys --> its 1st child as the new root if it has a child  
    // otherwise set root as NULL  
    if (root->n == 0) {  
        BTreeNode *tmp = root;  
        if (root->leaf)  
            root = NULL;  
        else  
            root = root->C[0];  
        delete tmp; // free old root  
    }  
    return;  
}
```

B-tree – C++ implementation

- BTreeNode methods
 - Creation of a node
 - Find key
 - Traverse
 - Search

```
BTreeNode::BTreeNode(int t1, bool leaf1) {
    // copy the given minimum degree and leaf property !
    t = t1; leaf = leaf1;
    // Allocate memory for maximum number of possible keys and child pointers
    keys = new int[2 * t - 1];
    C = new BTreeNode *[2 * t];
    n = 0; // Set the number of keys to 0
}

// returns the index of the 1st key >= k
int BTreeNode::findKey(int k) {
    int idx = 0;
    while (idx < n && keys[idx] < k)
        ++idx;
    return idx;
}
```

```
// traverse all the nodes in a subtree rooted with this node
void BTreeNode::traverse() {
    // There are n keys and n+1 children, traverse through n keys and 1st n children
    int i;
    for (i = 0; i < n; i++) {
        // if not leaf then before printing key[i]
        // traverse the subtree rooted with child C[i]
        if (leaf == false)
            C[i]->traverse();
        cout << " " << keys[i];
    }
    if (leaf == false)
        C[i]->traverse();
}

// search key k in subtree rooted with this node
BTreeNode *BTreeNode::search(int k) {
    int i = 0;
    while (i < n && k > keys[i])
        i++;
    if (keys[i] == k)
        return this;
    if (leaf == true)
        return NULL;
    return C[i]->search(k);
}
```

B-tree – C++ implementation

■ Remove

```
// remove the key k from the sub-tree rooted with this node
void BTreeNode::remove(int k) {
    int idx = findKey(k);
    if (idx < n && keys[idx] == k) { // k present in this node
        if (leaf)
            removeFromLeaf(idx);
        else
            removeFromNonLeaf(idx);
    }
    else {
        if (leaf) {
            cout << "The key " << k << " is does not exist.\n";
            return;
        }
        // The key to be removed is in the sub-tree rooted with this node
        // flag = whether the key is present in the sub-tree rooted with the last child of this node
        bool flag = ((idx == n) ? true : false);
        // If the child where the key is supposed to exist has less than t keys, fill that child
        if (C[idx]->n < t)
            fill(idx);
        // If the last child has been merged, must have merged with the previous child
        // we keep going on the idx-1 th child., otherwise we go on the idx th child (now has at least t keys)
        if (flag && idx > n)
            C[idx - 1]->remove(k);
        else
            C[idx]->remove(k);
    }
    return;
}
```

```
void BTreeNode::removeFromLeaf(int idx) {
    // Move all the keys after the idx-th pos one place backward
    for (int i = idx + 1; i < n; ++i)
        keys[i - 1] = keys[i];
    // Reduce the count of keys
    n--;
    return;
}

// A function to remove the idx-th key from this node - which is a non-leaf node
void BTreeNode::removeFromNonLeaf(int idx) {
    int k = keys[idx];
    // If the child that precedes k (C[idx]) has atleast t keys,
    // find the predecessor 'pred' of k in the subtree rooted at
    // C[idx]. Replace k by pred. Recursively delete pred
    // in C[idx]
    if (C[idx]->n >= t) {
        int pred = getPred(idx);
        keys[idx] = pred;
        C[idx]->remove(pred);
    }
    // If the child C[idx] has less than t keys, examine C[idx+1].
    // If C[idx+1] has atleast t keys, find the successor 'succ' of k in
    // the subtree rooted at C[idx+1]
    // Replace k by succ
    // Recursively delete succ in C[idx+1]
    else if (C[idx + 1]->n >= t) {
        int succ = getSucc(idx);
        keys[idx] = succ;
        C[idx + 1]->remove(succ);
    }
    // If both C[idx] and C[idx+1] has less than t keys, merge k and all of C[idx+1]
    // into C[idx]
    // Now C[idx] contains 2t-1 keys
    // Free C[idx+1] and recursively delete k from C[idx]
    else {
        merge(idx);
        C[idx]->remove(k);
    }
    return;
}
```

B-tree – C++ implementation

■ Remove

➤ Utility functions

```
// merge C[idx] with C[idx+1] C[idx+1] is freed after merging
void BTreeNode::merge(int idx) {
    BTreeNode *child = C[idx];
    BTreeNode *sibling = C[idx + 1];
    // Pulling a key from the current node and inserting it into (t-1)th position of C[idx]
    child->keys[t - 1] = keys[idx];
    // Copying the keys from C[idx+1] to C[idx] at the end
    for (int i = 0; i < sibling->n; ++i)
        child->keys[i + t] = sibling->keys[i];
    // Copying the child pointers from C[idx+1] to C[idx]
    if (!child->leaf) {
        for (int i = 0; i <= sibling->n; ++i)
            child->C[i + t] = sibling->C[i];
    }
    // Moving all keys after idx in the current node one step before -
    // to fill the gap created by moving keys[idx] to C[idx]
    for (int i = idx + 1; i < n; ++i)
        keys[i - 1] = keys[i];
    // Moving the child pointers after (idx+1) in the current node one step before
    for (int i = idx + 2; i <= n; ++i)
        C[i - 1] = C[i];
    // Updating the key count of child and the current node
    child->n += sibling->n + 1;
    n--;
    delete(sibling); // free memory occupied by sibling
    return;
}
```

```
// A function to get predecessor of keys[idx]
int BTreeNode::getPred(int idx) {
    // Keep moving to the right most node until we reach a leaf
    BTreeNode *cur = C[idx];
    while (!cur->leaf)
        cur = cur->C[cur->n];
    return cur->keys[cur->n - 1]; // the last key of the leaf
}

int BTreeNode::getSucc(int idx) {
    // Keep moving the left most node starting from C[idx+1] until we reach a leaf
    BTreeNode *cur = C[idx + 1];
    while (!cur->leaf)
        cur = cur->C[0];
    return cur->keys[0]; // the 1st key of the leaf
}

// fill child C[idx] which has less than t-1 keys
void BTreeNode::fill(int idx) {
    // If the previous child(C[idx-1]) more than t-1 keys --> borrow a key
    // from that child
    if (idx != 0 && C[idx - 1]->n >= t)
        borrowFromPrev(idx);
    // If the next child(C[idx+1]) more than t-1 keys --> borrow a key
    // from that child
    else if (idx != n && C[idx + 1]->n >= t)
        borrowFromNext(idx);
    // Merge C[idx] with its sibling
    else {
        if (idx != n) // If C[idx] == last child
            merge(idx); // merge with its previous sibling
        else
            merge(idx - 1); // merge with its next sibling
    }
    return;
}
```

B-tree – C++ implementation

■ Remove

➤ Utility functions: Borrow from **Previous** and from **Next**

```
// borrow a key from C[idx-1] and insert it into C[idx]
void BTreeNode::borrowFromPrev(int idx) {
    BTreeNode *child = C[idx];
    BTreeNode *sibling = C[idx - 1];
    // The last key from C[idx-1] goes up to the parent and key[idx-1]
    // from parent is inserted as the 1st key in C[idx]
    // --> the loses sibling one key and child gains one key
    // Moving all key in C[idx] one step ahead
    for (int i = child->n - 1; i >= 0; --i)
        child->keys[i + 1] = child->keys[i];
    // If C[idx] is not a leaf, move all its child pointers one step ahead
    if (!child->leaf) {
        for (int i = child->n; i >= 0; --i)
            child->C[i + 1] = child->C[i];
    }
    // Setting child's 1st key == to keys[idx-1] from the current node
    child->keys[0] = keys[idx - 1];
    // Moving sibling's last child as C[idx]'s 1st child
    if (!(child->leaf))
        child->C[0] = sibling->C[sibling->n];
    // Moving the key from the sibling to the parent
    // This reduces the number of keys in the sibling
    keys[idx - 1] = sibling->keys[sibling->n - 1];
    child->n += 1;
    sibling->n -= 1;
    return;
}
```

```
// borrow a key from the C[idx+1] and place it in C[idx]
void BTreeNode::borrowFromNext(int idx) {
    BTreeNode *child = C[idx];
    BTreeNode *sibling = C[idx + 1];
    // keys[idx] is inserted as the last key in C[idx]
    child->keys[(child->n)] = keys[idx];
    // Sibling's 1st child is inserted as the last child into C[idx]
    if (!(child->leaf))
        child->C[(child->n) + 1] = sibling->C[0];
    // 1st key from sibling is inserted into keys[idx]
    keys[idx] = sibling->keys[0];
    // Moving all keys in sibling one step behind
    for (int i = 1; i < sibling->n; ++i)
        sibling->keys[i - 1] = sibling->keys[i];
    // Moving the child pointers one step behind
    if (!sibling->leaf) {
        for (int i = 1; i <= sibling->n; ++i)
            sibling->C[i - 1] = sibling->C[i];
    }
    // Increasing and decreasing the key count of C[idx] and C[idx+1]
    child->n += 1;
    sibling->n -= 1;
    return;
}
```


B-tree – C++ implementation

■ Insert

```
// insert a new key in this node, the node must be non-full when insertNonFull is called
void BTreeNode::insertNonFull(int k) {
    // Initialize index as index of rightmost element
    int i = n - 1;
    // If this is a leaf node
    if (leaf == true) {
        // 1) Finds the location of new key to be inserted
        // 2) Moves all greater keys to 1 place ahead
        while (i >= 0 && keys[i] > k) {
            keys[i + 1] = keys[i];
            i--;
        }
        // Insert the new key at found location
        keys[i + 1] = k;
        n = n + 1;
    }
    else { // If this node is not leaf
        // Find the child which is going to have the new key
        while (i >= 0 && keys[i] > k)
            i--;
        // See if the found child is full
        if (C[i + 1] ->n == 2 * t - 1) {
            // If the child is full, then split it
            splitChild(i + 1, C[i + 1]);
            // After split, the middle key of C[i] goes up and
            // C[i] is splitted into 2
            // See which of the 2 is going to have the new key
            if (keys[i + 1] < k)
                i++;
        }
        C[i + 1] ->insertNonFull(k);
    }
}
```

```
// split the child y of this node (y must be full when splitChild is called)
void BTreeNode::splitChild(int i, BTreeNode *y) {
    // Create a new node which is going to store (t-1) keys of y
    BTreeNode *z = new BTreeNode(y ->t, y ->leaf);
    z ->n = t - 1;
    // Copy the last (t-1) keys of y to z
    for (int j = 0; j < t - 1; j++)
        z ->keys[j] = y ->keys[j + t];
    // Copy the last t children of y to z
    if (y ->leaf == false) {
        for (int j = 0; j < t; j++)
            z ->C[j] = y ->C[j + t];
    }
    // Reduce the number of keys in y
    y ->n = t - 1;
    // As this node is going to have a new child, create space of new child
    for (int j = n; j >= i + 1; j--)
        C[j + 1] = C[j];
    // Link the new child to this node
    C[i + 1] = z;
    // A key of y will move to this node. Find location of
    // new key and move all greater keys one space ahead
    for (int j = n - 1; j >= i; j--)
        keys[j + 1] = keys[j];
    // Copy the middle key of y to this node
    keys[i] = y ->keys[t - 1];
    // Increment count of keys in this node
    n = n + 1;
}
```


Conclusion

■ B-trees

➤ Efficient data structure

- Keep keys in sorted order for sequential traversing
- Use a hierarchical index to minimize the number of disk reads
- Use partially full blocks to speed insertions and deletions
- Keep the index balanced with a recursive algorithm

■ Complexity

➤ 2-3 tree & B-tree

Algorithm	Average	Worst case
Space	$O(n)$	$O(n)$
Search	$O(\log n)$	$O(\log n)$
Insert	$O(\log n)$	$O(\log n)$
Delete	$O(\log n)$	$O(\log n)$

Questions ?

- Acknowledgment + Reading
 - Chapter 18, B-trees, Introduction to Algorithms, 3rd Edition.

