


# Algorithms and Data Structures (CSci 115)

California State University Fresno  
College of Science and Mathematics  
Department of Computer Science  
H. Cecotti

# Learning outcomes

## ■ Red-Black trees

- Definitions
- How to search, insert, remove elements



# Definitions

- Red-Black Tree (RBT) (1972 – Rudolph Bayer)
  - Binary Search Tree (BST)
    - with one extra **bit** of storage per node:
      - its color: **RED** or **BLACK**.
  - Constraint of the node colors on any simple path from the root to a leaf,
    - RBT ensure that:
      - no such path is more than 2 times as long as any other,
      - → the tree is approximately balanced.
  - Each node of the tree contains the attributes
    1. Color
    2. Key
    3. Left
    4. Right
    5. Parent
  - If a child or the parent of a node does not exist
    - the corresponding pointer attribute of the node contains the value NIL. (NIL == NULL == null pointer)
  - NULL == pointers to leaves (external nodes) of the BST
  - the key-bearing nodes == internal nodes of the tree.

# Properties


---

- 5 key properties
  1. Every node is either **red** or **black**.
  2. The root is **black**.
  3. Every leaf (NIL) is **black**.
  4. If (a node is **red**) then
    - **both** its children are **black**.
  5. For each node, all simple paths from the node to descendant leaves contain the **same** number of **black** nodes.

# Example

- RBT


- with leaves and the root's parent omitted entirely.



# Example

## ■ RBT

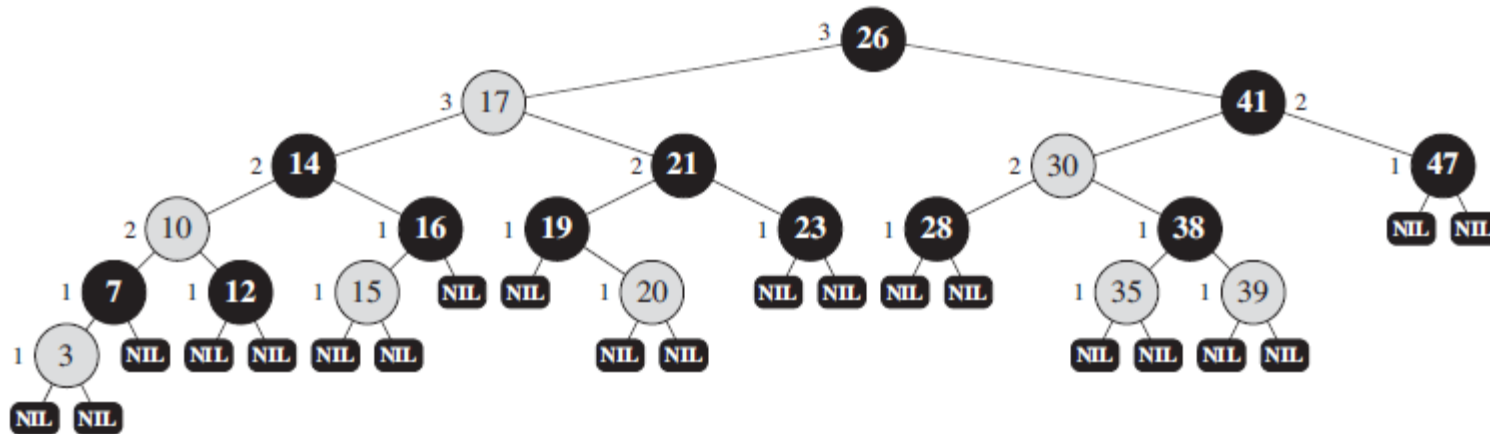
- with each NIL replaced by the single sentinel  $T.nil$ ,
  - which is always **black**, and with black-heights omitted.



# Example

## ■ RBT

- Every leaf, shown as a NIL = **black**.
- Each non-NIL node is marked with its **black-height**
- NILs have black-height 0.



# Definitions

---

- The **black-height** of the node:  $bh(x)$ 
  - Number of **black** nodes on any simple path from, but not including, a node  $x$  down to a leaf
- Property 5 (red  $\rightarrow$  both children are black)  $\rightarrow$ 
  - well defined notion of **black-height** because
    - all descending simple paths from the node have the **same number of black nodes**.
  - We define the **black-height** of a RBT == the **black-height** of its **root**.
- RBT with  $n$  internal nodes  $\rightarrow$  height at most  $2 * \log(n+1)$



# Proof (1)

- Let  $h$  be the height of the tree.
- Show that
  - Subtree rooted at any node  $x$  contains at least  $2^{bh(x)}-1$  internal nodes.
- By induction on the height of  $x$ .
  - If  $(h(x)=0)$  then
    - $x$  must be a leaf ( $T.nil$ )
    - the subtree rooted at  $x$  contains at least  $2^{bh(x)}-1 = 0$  internal nodes.
  - Inductive step
    - consider a node  $x$  with  $h>0$  and  $x$  is an internal node with 2 children.
    - Each child has  $bh(x)$  (**red**) or  $bh(x)-1$  (**black**)
      - Because the  $h(x.child) < h(x)$
  - Apply the inductive hypothesis
  - → each child has at least  $2^{bh(x)-1}-1$  internal nodes
  - → the subtree rooted at  $x$  contains **at least**
    - $(2^{bh(x)-1}-1)+(2^{bh(x)-1}-1) + 1 = 2^{bh(x)}-1$  internal nodes

# Proof (2)

---

- Property 4
  - At least half the nodes on any simple path from the root to a leaf, not including the root, must be **black**
- $\rightarrow bh(\text{root})$  must be at least  $h/2$
- $\rightarrow n \geq 2^{h/2} - 1$ 
  - $\rightarrow \log(n+1) \geq h/2$
  - $2 \log(n+1) \geq h$
- We can implement Search, Minimum, Maximum, Successor, and Predecessor in  $O(\log n)$

# Relationships with B-trees

## ■ RBT

### ➤ Same structure as a B-tree of order 4

- Each node can contain
  - between 1 and 3 values
  - between 2 and 4 child pointers

### ➤ In this B-tree, each node will contain only 1 value matching the value in a **black** node of the RBT

- with an optional value before and/or after it in the same node
- matching an equivalent **red** node of the RBT

## ■ Move up the **red** nodes

### ➤ → they align horizontally with their parent **black** node


- Through the creation of an horizontal cluster.
- B-tree, / modified graphical representation of the RBT → all leaf nodes have the same depth.

## ■ A minimum fill factor of 33% of values per cluster


### ➤ with a maximum capacity of 3 values.

# Relationships with B-trees

- RBT versus B-Tree
  - Property 4



RBT




B-tree

# Red Black Tree

## ■ Insert & Delete

- We must keep the constraints of the tree
- → change the colors of some of the nodes in the RBT
- → change the pointer structure of the RBT
  - Through Rotation
    - Like AVL trees



# Insertion & Deletion

---

- Node z
  - Parent: z.p
  - Grand-parent: z.p.p
  - Uncle (y)
    - z.p.p.left
      - if z.p.p.right == z.p
    - z.p.p.right
      - if z.p.p.left == z.p

# Insertion

- Insert node **z** in tree **T**
  - Default insert like in a BST
  - Set color to RED
  - InsertFixup

Find where to insert  
Same as BST

Special case: root?

Default color: **RED**

RB-INSERT(*T*, *z*)

```
1  y = T.nil
2  x = T.root
3  while x ≠ T.nil
4      y = x
5      if z.key < x.key
6          x = x.left
7      else x = x.right
8  z.p = y
9  if y == T.nil
10     T.root = z
11  elseif z.key < y.key
12     y.left = z
13  else y.right = z
14  z.left = T.nil
15  z.right = T.nil
16  z.color = RED
17  RB-INSERT-FIXUP(T, z)
```

© 2006 The Authors  
Journal compilation © 2006 Blackwell Publishing Ltd



# Insertion

## ■ Fix tree

### ➤ Pseudo Code

RB-INSERT-FIXUP( $T, z$ )

```
1  while  $z.p.color == RED$ 
2      if  $z.p == z.p.p.left$ 
3           $y = z.p.p.right$      $z.p.p = \text{grand parent}$ 
4          if  $y.color == RED$    $y \text{ is RED}$ 
5               $z.p.color = BLACK$     // case 1
6               $y.color = BLACK$       // case 1
7               $z.p.p.color = RED$     // case 1
8               $z = z.p.p$             // case 1
9          else if  $z == z.p.right$    $y \text{ is Black}$ 
10              $z = z.p$               // case 2
11             LEFT-ROTATE( $T, z$ )    // case 2
12              $z.p.color = BLACK$     // case 3
13              $z.p.p.color = RED$     // case 3
14             RIGHT-ROTATE( $T, z.p.p$ ) // case 3
15         else (same as then clause
              with "right" and "left" exchanged)
16      $T.root.color = BLACK$ 
```



What we want to keep in the loop (**invariant**)

1. Node  $z$  is red
2. If ( $z.p$  is the root) then  $z.p$  is black
3. If the tree violates any of the RBT properties then it violates at most one of them and the violation is property 2 or property 4.
  - If property 2 violation
  - $\rightarrow$  because  $z$  is the root and is red.
  - If property 4 violation
  - $\rightarrow$  because both  $z$  and  $z.p$  are red.


# Insertion

---

- **Case 1:**
  - z's uncle y is red
- **Case 2:**
  - z's uncle y is black and z is a right child
- **Case 3:**
  - z's uncle y is black and z is a left child

# Insert - Case 1


- Violation of Property 4 (if a node is red, both children are black)
  - $z$  and  $z.p$  : red



# Insert - Case 2 & 3

## ■ Violation of Property 4

➤ z and z.p: both red ☹️



Each of the subtrees  $\alpha$ ,  $\beta$ , and  $\gamma$  has a black root

- $\alpha$ ,  $\beta$ , and  $\gamma$  from property 4
- $\delta$  because case 1 otherwise

each has the same bh (black height)

Transform Case 2  $\rightarrow$  Case 3 by a **left rotation** (preserves property 5)

- all downward simple paths from a node to a leaf have the same number of blacks.
- Case 3 causes some **color changes** and a **right rotation** (preserve property 5).

The **while** loop then terminates because property 4 is satisfied == no longer 2 red nodes in a row.

# Insert

---

## ■ Complexity

- Height of a RBT on  $n$  nodes =  $O(\log n)$
- → RB-INSERT take  $O(\log n)$  in time.
- RB-INSERTFIXUP,
  - **while** loop repeats **only if**
    - case 1 happens and then the pointer  $z$  moves 2 levels up the tree.
    - → total number of times the **while** loop can be executed is  $O(\log n)$
  - →  $O(\log n)$  time.
- It never performs more than 2 rotations
  - because since the **while** loop terminates if case 2 or case 3 is executed!!

# Delete

## ■ Transplant function

- It replaces one subtree as a child of its parent with another subtree
  - replaces the subtree rooted at node **u** with the subtree rooted at node **v**
- /!\ Reference to **T.nil** , not **nil**
  - Use of a sentinel

RB-TRANSPLANT(*T*, *u*, *v*)

```
1  if u.p == T.nil
2      T.root = v
3  elseif u == u.p.left
4      u.p.left = v
5  else u.p.right = v
6  v.p = u.p
```

## ■ Difference with BST delete

- Keep track of potential issues for violation of the RBT properties

# Delete

---

## ■ Main idea

### ➤ Delete node z

- If (z has fewer than 2 children) then
  - z is removed from the tree
- $y = z$
- If (z has 2 children) then
  - y should be z's successor
  - y moves into z's position in the tree.
  - Remember y's color before it is removed from or moved within the tree
  - Keep track of the node x that moves into y's original position in the tree
    - since node x might also cause violations
- After deleting node z,
  - RB-DELETE calls RB-DELETE-FIXUP
    - changes colors and performs rotations to restore the RBT properties.

# Delete

## ■ Pseudo code

➤  $O(\log(n))$

RB-Transplant( $T, u, v$ ):

replaces the subtree rooted at node  $u$  with the subtree rooted at node  $v$

RB-DELETE( $T, z$ )

```
1   $y = z$ 
2   $y\text{-original-color} = y.\text{color}$ 
3  if  $z.\text{left} == T.\text{nil}$ 
4       $x = z.\text{right}$ 
5      RB-TRANSPLANT( $T, z, z.\text{right}$ )
6  elseif  $z.\text{right} == T.\text{nil}$ 
7       $x = z.\text{left}$ 
8      RB-TRANSPLANT( $T, z, z.\text{left}$ )
9  else  $y = \text{TREE-MINIMUM}(z.\text{right})$ 
10      $y\text{-original-color} = y.\text{color}$ 
11      $x = y.\text{right}$ 
12     if  $y.p == z$ 
13          $x.p = y$ 
14     else RB-TRANSPLANT( $T, y, y.\text{right}$ )
15          $y.\text{right} = z.\text{right}$ 
16          $y.\text{right}.p = y$ 
17     RB-TRANSPLANT( $T, z, y$ )
18      $y.\text{left} = z.\text{left}$ 
19      $y.\text{left}.p = y$ 
20      $y.\text{color} = z.\text{color}$ 
21 if  $y\text{-original-color} == \text{BLACK}$ 
22     RB-DELETE-FIXUP( $T, x$ )
```



# Delete

---

- If ( $y.c == \text{red}$ )
  - RBT properties are respected when  $y$  is removed because:
    1. No bh in the tree have changed.
    2. No red nodes have been made adjacent.
      - Because  $y$  takes  $z$ 's place in the tree, along with  $z$ 's color
      - $\rightarrow$  cannot have 2 adjacent red nodes at  $y$ 's new position in the tree.
      - if  $y$  was not  $z$ 's right child then  $y$ 's original right child  $x$  replaces  $y$  in the tree
      - If ( $y.c == \text{red}$ ) then  $x.c == \text{black}$ 
        - so replacing  $y$  by  $x$  cannot cause 2 red nodes to become adjacent!
    3. As  $y$  could not have been the root if it was red
      - $\rightarrow$  root remains black.

# Delete

- If (y.c==black)
  - Remove y → change hb
  - Problems to fix the tree ! ☹️
- 1. If y had been the root and a red child of y becomes the new root
  - → violation of property 2.
- 2. If (x.c==red) and (x.p.c==red)
  - → violation of property 4.
- 3. Moving y within the tree causes any simple path that previously contained y to have 1 fewer black node
  - → violation of property 5 (by any ancestor of y in the tree!)
  - Correction
    - say that node x, now occupying y's original position, has an "extra" black
    - if (add 1 to the count of black nodes on any simple path that contains x)
    - Then under this interpretation, property 5 holds. 🤔
  - When we remove or move the black node y → we "push" its blackness onto node x.
- Problem:
- now node x is neither red nor black → violation of property 1.

# Delete

- Delete Fixup
  - Pseudo-code




RB-DELETE-FIXUP( $T, x$ )

```
1  while  $x \neq T.root$  and  $x.color == BLACK$ 
2      if  $x == x.p.left$ 
3           $w = x.p.right$ 
4          if  $w.color == RED$ 
5               $w.color = BLACK$  // case 1
6               $x.p.color = RED$  // case 1
7              LEFT-ROTATE( $T, x.p$ ) // case 1
8               $w = x.p.right$  // case 1
9          if  $w.left.color == BLACK$  and  $w.right.color == BLACK$ 
10              $w.color = RED$  // case 2
11              $x = x.p$  // case 2
12         else if  $w.right.color == BLACK$ 
13              $w.left.color = BLACK$  // case 3
14              $w.color = RED$  // case 3
15             RIGHT-ROTATE( $T, w$ ) // case 3
16              $w = x.p.right$  // case 3
17              $w.color = x.p.color$  // case 4
18              $x.p.color = BLACK$  // case 4
19              $w.right.color = BLACK$  // case 4
20             LEFT-ROTATE( $T, x.p$ ) // case 4
21              $x = T.root$  // case 4
22         else (same as then clause with “right” and “left” exchanged)
23              $x.color = BLACK$ 
```

# Delete (Fixup)

## ■ Case 1:

- It happens when
  - node  $w$  (sibling of  $x$  is red)  $w.c == \text{red}$
- $w$  must have black children
  - we can **switch** the colors of  $w$  and  $x.p$
  - Perform a **left-rotation** on  $x.p$  without violating any RBT properties
- New sibling of  $x$  (1 of  $w$ 's children prior to the rotation): **black**
- case 1  $\rightarrow$  case 2, 3, or 4.



# Delete (Fixup)

## ■ Case 2:

➤ It happens when

- (x's sibling w.c==black) and (w.right.c==black) and (w.left.c==black)

➤ As w is black → take one black off both x and w

- leaving x with only 1 black and leaving w red.


➤ Compensate for removing 1 black from x and w

- Add an extra black to x.p that was red or black
- How?

- by repeating the **while** loop with x.p as the new node x.
- If we enter **Case 2** through **Case 1**, the new node x is red-and-black, since the original x.p was red.


■ → color of the new node x is red

➤ and the loop terminates when it tests the loop condition.



# Delete (Fixup)

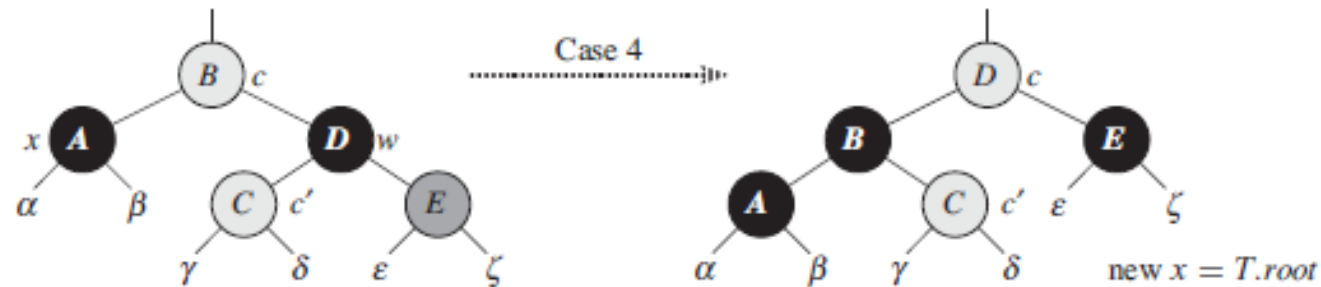
- **Case 3:**
  - It happens when
    - $w.c == \text{black}$  and  $w.\text{left}.c == \text{red}$  and  $w.\text{right}.c == \text{black}$
  - Switch the colors of  $w$  and  $w.\text{left}$
  - Then perform a **right rotation** on  $w$  without
    - violating any of the RBT properties
- The new sibling  $w$  of  $x$  is now a **black** node
  - with a red right child
  - Transform Case 3  $\rightarrow$  case 4.



# Delete (Fixup)

## ■ Case 4:

- It happens when
  - (x's sibling w.c==black) and (w.right.c==red)
- Some color changes + left rotation on x.p
  - Remove the extra black on x, making it singly black
    - without violating any of the RBT properties.
- Setting x to be the root causes the while loop to terminate when it tests the loop condition.



# Delete

## ■ Complexity analysis

➤ height of a RBT of  $n$  is  $O(\log n)$ ,

- total cost of the procedure without the call to RB-DELETEFIXUP is  $\log n$  time

➤ Within RB-DELETE-FIXUP

- Cases 1, 3, and 4:

- lead to termination after performing a **constant number of color changes** and
- + At most 3 rotations

- Case 2 is the **only** case where the **while** loop can be repeated

- the pointer  $x$  moves **up** the tree at most  $O(\log n)$  times,
- no rotations!

➤ → RB-DELETE-FIXUP

- takes  $O(\log n)$  time
- Performs at most 3 rotations

➤ Total time for RB-DELETE =  $O(\log n)$



# Conclusion

## ■ Red-Black trees

- A form of semi-balanced binary tree
- Compared to other trees
  - more general purpose: add, remove, and look-up
  - but AVL trees have faster look-ups at the cost of slower add/remove !

## ■ Applications

### ➤ Linux kernel

- The anticipatory, deadline, and CFQ I/O schedulers: RBT to track requests
- The packet CD/DVD driver: RBT trees
- The high-resolution timer uses an RBT to organize outstanding timer requests
- The ext3 file system tracks directory entries in an RBT
- Virtual memory areas (VMAs) are tracked with RBTs



## ■ Complexity

Algorithm	Average	Worst case
Space	$O(n)$	$O(n)$
Search	$O(\log n)$	$O(\log n)$
Insert	$O(\log n)$	$O(\log n)$
Delete	$O(\log n)$	$O(\log n)$

# Questions ?

---

- Acknowledgment + Reading
  - Csci 115 book Section 7.4
  - Chapter 13, Red-Black Trees, Introduction to Algorithms, 3<sup>rd</sup> Edition.

