

Algorithms and Data Structures (CSci 115)

California State University Fresno
College of Science and Mathematics
Department of Computer Science
H. Cecotti

Learning outcomes

- Multithreading*
 - Example with the sequence of Fibonacci
- Memory Optimization for data structures*

* : Not in the final exam

Introduction

- Link between **software** and **hardware**

- Everything is connected

- Importance of considering a **holistic** (global) approach

- → Direct link with **Operating Systems (OS)**

- 1 vs. Multi threads
 - Shared vs. Distributed memory

- → Think about

- **where** the application will be deployed
 - How often it will be used, how often the different actions will be used

Rationale

■ Considerations

➤ Serial algorithms

- For running on a uniprocessor computer
 - Only 1 instruction executes at a time

➤ Parallel algorithms

- To run on a multiprocessor computer that permits multiple
 - → instructions to execute concurrently

➤ Parallel computer

- **Shared memory**
 - Each processor can directly access any location of memory
- **Distributed memory**
 - Each processor's memory is private and an explicit message must be sent between processors in order for one processor to access the memory of another

Concurrency keywords

▪ Concurrency keywords

➤ Spawn

- **If** (spawn proceeds a procedure call)
- **then** the procedure instance that executes the spawn (the parent) may continue to execute in parallel with the spawned subroutine (the child), instead of waiting for the child to complete.
- The keyword spawn does *not* say that a procedure must execute concurrently, but simply that **it may**.
- At runtime:
 - It is up to the scheduler to decide which subcomputations should run concurrently.


➤ Sync

- The procedure must wait as necessary for all its spawned children to complete before proceeding to the statement after the sync

Back to Fibonacci (again)


■ Fibonacci

- *Golden ratio: 1.61803*
- Sequence: 0, 1, 1, 2, 3, 5, 8, 13, 21, 34, ...
- One of the first examples given for recursive functions
- Dynamic programming
 - Memoization
- Fibonacci heaps



```
FIB(n)
1  if  $n \leq 1$ 
2      return  $n$ 
3  else  $x = \text{FIB}(n - 1)$ 
4       $y = \text{FIB}(n - 2)$ 
5      return  $x + y$ 
```

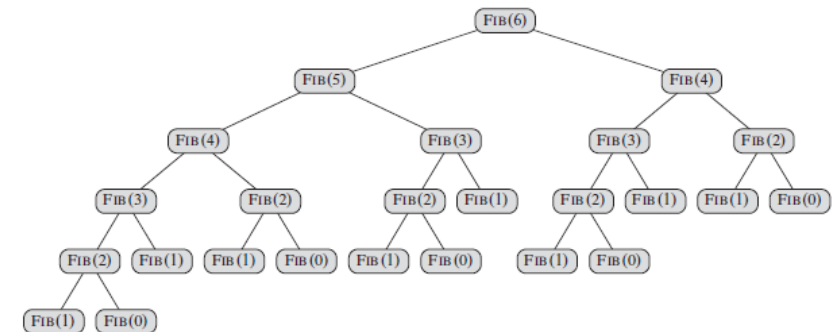
■ Sequence:



Back to Fibonacci (again)

- Exploration of dynamic multithreading
 - With the sequence of Fibonacci
- Let $T(n)$: the running time of $\text{Fibonacci}(n)$
 - Since this procedure contains 2 recursive calls and a constant amount of extra work
 - $T(n) = T(n-1) + T(n-2) + \theta(1)$
 - $\rightarrow T(n) = \theta(F_n) = \theta\left(\left(\frac{1+\sqrt{5}}{2}\right)^n\right)$
 - **Exponential growth**
 - Particularly bad way to calculate Fibonacci numbers.
 - How would you calculate the Fibonacci numbers?

```
FIB(n)
1  if  $n \leq 1$ 
2      return  $n$ 
3  else  $x = \text{FIB}(n-1)$ 
4       $y = \text{FIB}(n-2)$ 
5      return  $x + y$ 
```



Back to Fibonacci

- Implementation with matrixes

$$\begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}^n = \begin{pmatrix} F_{n+1} & F_n \\ F_n & F_{n-1} \end{pmatrix}$$

- To calculate F_n in $O(\log n)$ steps
 - by repeated squaring of the matrix
 - you can calculate the Fibonacci numbers with a **serial** algorithm.
- To illustrate the principles of parallel programming
 - Use the naive (bad) algorithm

Fibonacci Example

- Parallel algorithm to compute Fibonacci numbers:
 - Fibonacci(n)
 - *if* ($n < 2$) *then* return n ;
 - $x = \text{spawn Fibonacci}(n-1)$; // parallel execution
 - $y = \text{spawn Fibonacci}(n-2)$; // parallel execution
 - **sync**; // wait for results of x and y
 - return $x + y$;

Computation DAG

- Multi-threaded computation
 - Using the help of a computation **Directed Acyclic Graph (DAG)** $G=(V,E)$.
 - V : instructions.
 - E : dependencies between instructions.
- An edge (u,v) is in E
 - the instruction u must execute **before** instruction v .
- A computation DAG $G=(V,E)$
 - It consists of :
 - The vertex set V : the threads of the program.
 - The edge set E contains an edge (u,v) if and only if the thread u need to execute before thread v .
 - **If** $(\text{ExistEdge}(u,v))$
 - **then** they are said to be (logically) in **series**
 - **If** (there is no thread)
 - **then** they are said to be (logically) in **parallel**

Strand and Threads

- **Definitions**

- A sequence of instructions containing **no parallel control** (spawn, sync, return from spawn, parallel) can be grouped into a **single strand**.
- A strand of maximal length: a thread.

Edge Classification

- A continuation edge (u,v) connects a thread u to its successor v within the same procedure instance.
- **If** (a thread u spawns a new thread v)
 - **Then** (u,v) is called a **spawn** edge.
- **If** (a thread v returns to its calling procedure **and** x is the thread following the parallel control)
 - **Then** the return edge (v,x) is included in the graph.

Fibonacci Example

- Parallel algorithm to compute Fibonacci numbers:

- Fibonacci(n)

- if $n < 2$ then return n ; // thread A

- $x = \text{spawn Fibonacci}(n-1);$

- $y = \text{spawn Fibonacci}(n-2);$ // thread B

- sync;

- return $x + y$; // thread C

Performance Measures

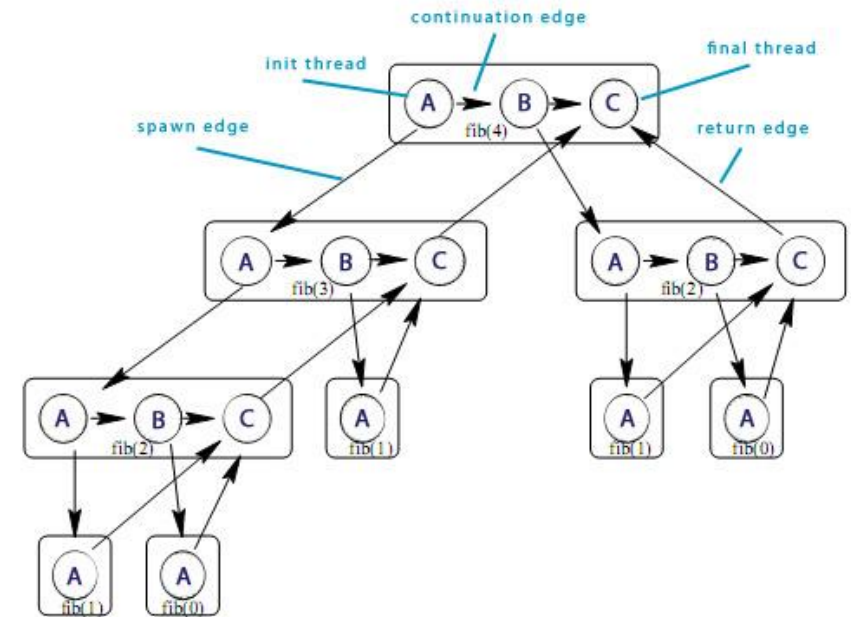
■ Definitions

➤ Work of a multithreaded computation:

- the total time to execute the entire computation on 1 processor.
- Work
 - sum of the times taken by each thread


➤ The span

- longest time to execute the threads along any path of the computational DAG



Performance Measures

- In Fibonacci(4)
 - 17 vertices = 17 threads.
 - 8 vertices on longest path.
- Assuming
 - unit time for each thread
- We obtain:
 - Work = 17 time units
 - Span = 8 time units



Performance Measures

- The actual running time of a multithreaded computation depends on:
 - its **work** and **span**
 - on how many processors/cores are available
 - how the scheduler allocates strands to processors
- Running time on P processors is indicated by subscript P
 - T_1 running time on a single processor
 - T_P running time on P processors
 - T_∞ running time on unlimited processors

Definitions

■ Work law

- An ideal parallel computer with P processors can do at most P units of work.
- Total work to do is $T_1 \rightarrow PT_p \geq T_1$
- The work law: $T_p \geq T_1/P$

■ Span law

- A P -processor ideal parallel computer cannot run faster than a machine with unlimited number of processors.
- However
 - A computer with unlimited number of processors can emulate a P -processor machine
 - by using simply P of its processors
 - \rightarrow The span law is:
 - $T_p \geq T_\infty$

- The speed up of a computation on P processors is defined as T_1 / T_p
- The parallelism of a multithreaded computation is given by T_1 / T_∞

Scheduling

- The performance depends not just on the work and span
- In addition,
 - the strands must be
 - scheduled efficiently.
 - mapped to static threads
 - and the OS schedules the threads on the processors themselves.
- The scheduler must schedule the computation
 - with no advance knowledge of when the strands will be spawned or when they will complete; it must operate online.

Greedy Scheduler

- We will assume a greedy scheduler in our analysis, since this keeps things simple. A greedy scheduler assigns as many strands to processors as possible in each time step.
- On P processors, if at least P strands are ready to execute during a time step, then we say that the step is a complete step; otherwise we say that it is an incomplete step.

Greedy Scheduler Theorem

- On an ideal parallel computer with P processors, a greedy scheduler executes a multithreaded computation with work T_1 and span T_∞ in time
 - $T_P \leq T_1 / P + T_\infty$
- As the best we can hope for on P processors is:
 - $T_P = T_1 / P$ by the work law
 - $T_P = T_\infty$ by the span law
 - the sum of these 2 lower bounds

Proof (part 1)

- Let's consider the complete steps
 - In each complete step, the P processors perform a total of P work.
- Seeking a **contradiction**
 - we assume that the number of complete steps exceeds T_1/P
 - then the total work of the complete steps is at least

$$\begin{aligned} P(\lfloor T_1/P \rfloor + 1) &= P\lfloor T_1/P \rfloor + P \\ &= T_1 - (T_1 \bmod P) + P \\ &> T_1 \end{aligned}$$

- As this exceeds the total work required by the computation
 - This is impossible 😊

Proof (part 2)

- Now consider an incomplete step.
 - Let G be the DAG representing the entire computation.
 - Without loss of generality, we assume that each strand takes unit time
 - Otherwise replace longer strands by a chain of unit-time strands!
- Let G' be the subgraph of G
 - that has yet to be executed at the **start** of the incomplete step
- Let G'' be the subgraph
 - remaining to be executed **after** the completion of the incomplete step.

Proof (part 3)

- A longest path in a DAG must necessarily start at a vertex with in-degree 0.
- Since an incomplete step of a greedy scheduler executes all strands with in-degree 0 in G' ,
 - the length of the longest path in G'' must be **1 less than the length of the longest path in G'** .
- Idea
 - An incomplete step **decreases** the span of the unexecuted DAG by **1**.
 - → the number of incomplete steps is at most T_∞
 - Since each step is either complete or incomplete, the theorem follows.

Corollary

- The running time of any multithreaded computation scheduled
 - by a greedy scheduler on an ideal parallel computer with P processors
- is **within a factor of 2 of optimal**.
- Proof:
 - The T_P^* be the running time produced by an optimal scheduler
 - Let T_1 be the work and T_∞ be the span of the computation
 - Then $T_P^* \geq \max(T_1 / P, T_\infty)$
 - By the theorem,
 - $T_P \leq T_1 / P + T_\infty \leq 2 \max(T_1 / P, T_\infty) \leq 2 T_P^*$

Slackness

▪ Definitions

- The parallel slackness of a multithreaded computation executed on an ideal parallel computer with P processors is the ratio of parallelism by P .
- $\text{Slackness} = (T_1 / T_\infty) / P$
- **If** (the slackness is less than 1)
- **then** we cannot hope to achieve a linear speedup.

Speedup

- Let T_P be the running time of a multi-threaded computation
 - produced by a greedy scheduler on an ideal computer with P processors
- Let T_1 be the work and T_∞ be the span of the computation
- If the slackness is big, $P \ll (T_1 / T_\infty)$
- then T_P is approximately T_1 / P .
- Proof:
 - If $P \ll (T_1 / T_\infty)$
 - then $T_\infty \ll T_1 / P$
 - \rightarrow By the theorem, we have
 - $T_P \leq T_1 / P + T_\infty \approx T_1 / P$
 - By the work law, we have
 - $T_P \geq T_1 / P$
 - $\rightarrow T_P \approx T_1 / P$, as claimed.

Work of Fibonacci

- We want to know the **work** and **span** of the Fibonacci computation
 - to compute the parallelism (work/span) of the computation.
- The work T_1 is straight forward
 - since it amounts to compute the running time of the serialized algorithm.
- $T_1 = \theta(((1+\text{sqrt}(5))/2)^n)$

Span of Fibonacci

- Recall that the span T_∞ is the longest path in the computational DAG
 - Since $\text{Fibonacci}(n)$ spawns
 - $\text{Fibonacci}(n-1)$
 - $\text{Fibonacci}(n-2)$
 - We have
 - $T_\infty(n) = \max(T_\infty(n-1) , T_\infty(n-2)) + \theta(1) = T_\infty(n-1) + \theta(1)$
 - which yields $T_\infty(n) = \theta(n)$.

Parallelism of Fibonacci

- The parallelism of the Fibonacci computation is
 - $T_1(n)/T_\infty(n) = \theta(((1+\text{sqrt}(5))/2)^n / n)$
 - which grows dramatically as n gets large.
- → Even on the largest parallel computers
 - A modest value of n suffices to achieve near perfect linear speedup

Data structure

■ The idea

- Keep heavily accessed data members near each other physically in memory a system's caching
- If it's declared together, it's easier to access it together
 - In the same function

■ Organization:

- Field reordering
 - In a struct with multiple elements,
 - keep the items that will be accessed together, together in the order of elements in the struct

Software prefetching

■ Definition

➤ Cache prefetching:

- Technique used by computer processors to boost execution performance
- How?: by fetching instructions/data
 - **from** their original storage in “**slow memory**”
 - **to** a “**faster local memory**” before it is actually needed

■ Approaches

➤ Not too early

- Data may be evicted before use

➤ Not too late

- Data not fetched in time for use

➤ Greedy

Conclusion

- To take advantage of the architecture:
 - Multi-threaded/multi processors/cores architectures
 - Memory available
 - Disk size available
- **Remark**
 - What is your job? What is the job of the compiler?
 - Modern compilers deal with many optimization
 - Your priority:
 - Depends on the project:
 - Team work, Maintenance, Code easy to read, update, maintain
 - Optimization...
- **Next session:**
 - Final conclusion of the CSc115 course + Revisions for the final

Questions ?

- Reading & Acknowledgement
 - Multithreaded Algorithms ,Introduction to Algorithms, 3rd Edition, Chapter 26.

